



Scientific Collaboration - GENCI, IBM, NVIDIA, Mellanox

Introduction To Ouessant OpenPOWER Platform

Stéphane CHAUX NVIDIA

schauveau@nvidia.com

François COURTEILLE NVIDIA

fcourteille@nvidia.com

Pascal VEZOLLE IBM

vezolle@fr.ibm.com

Nicolas TALLET IBM

nicolas.tallet@fr.ibm.com

OpenPOWER Scientific Collaboration - GENCI & IBM/NVIDIA/Mellanox

▪ Objectives of the Collaboration

- OpenPower Technology Assessment for Multi-Petaflops Deployment in 2017
 - Analyze applications affinity and requirements
 - Port applications to OpenPower platform
 - Analyze applications performance
 - Organize workshops and exchanges with application developers
- Experiment OpenPower w/GPU Prototype Systems in 2015/2016
- Evaluate Programming Models (OpenMP, OpenACC)
- Provide Feedback To Development Teams

OpenPOWER Technical Architecture

Hardware Features

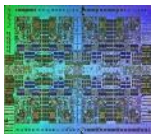
IBM OpenPOWER Accelerated Computing Roadmap

Mellanox
Interconnect
TechnologyConnectX-4
EDR Infiniband
PCIe Gen3ConnectX-4
EDR Infiniband
CAPI over PCIe Gen3ConnectX-5
HDR Infiniband
Enhanced CAPI over PCIe Gen4

NVIDIA GPUs

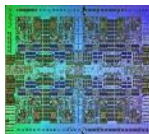
Kepler
PCIe Gen3Pascal
NVLink
SXM2Volta
Enhanced NVLink
SXM2

POWER8



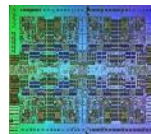
IBM CPUs

POWER8 w/NVLink



NVLink

POWER9

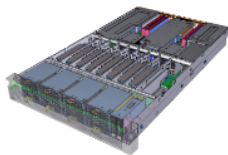
Enhanced
NVLink

2015

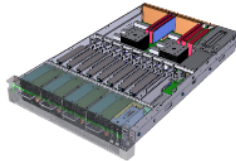
2016

2017

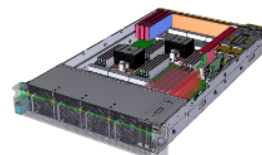
Server



Firestone



Minsky



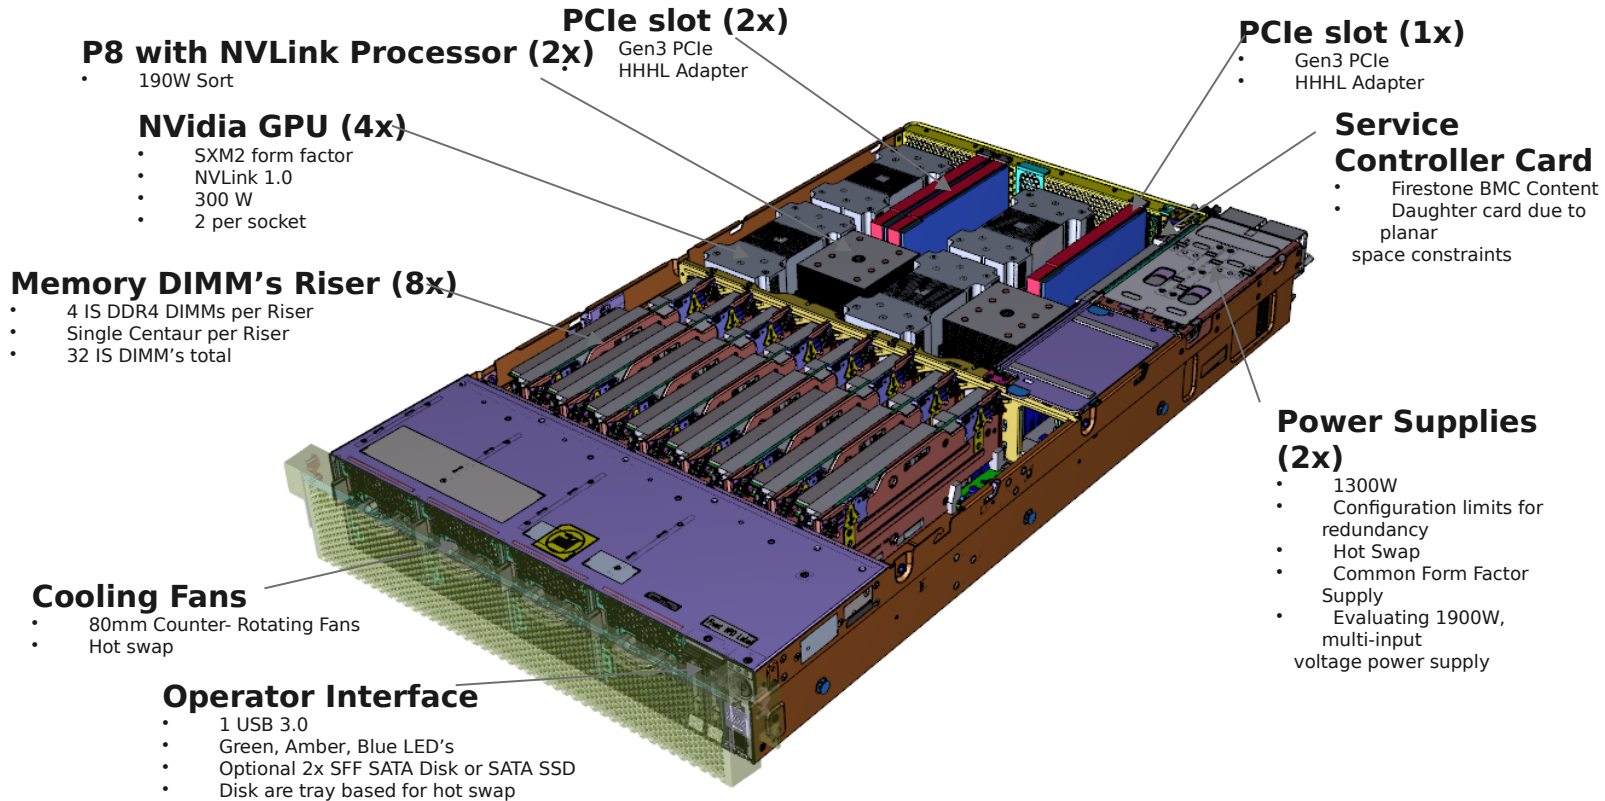
Witherspoon

IBM Power Systems S822LC 'Minsky' w/NVIDIA Tesla P100

- **2 POWER8 CPUs**
- **Up to 1TB DDR4 memory**
- **Up to 4 Tesla P100 GPUs**
- **1st Server with POWER8 with NVLink Technology**
- **Only architecture with CPU:GPU NVlink**



IBM Power System S822LC 'Minsky' (8335-GTB)



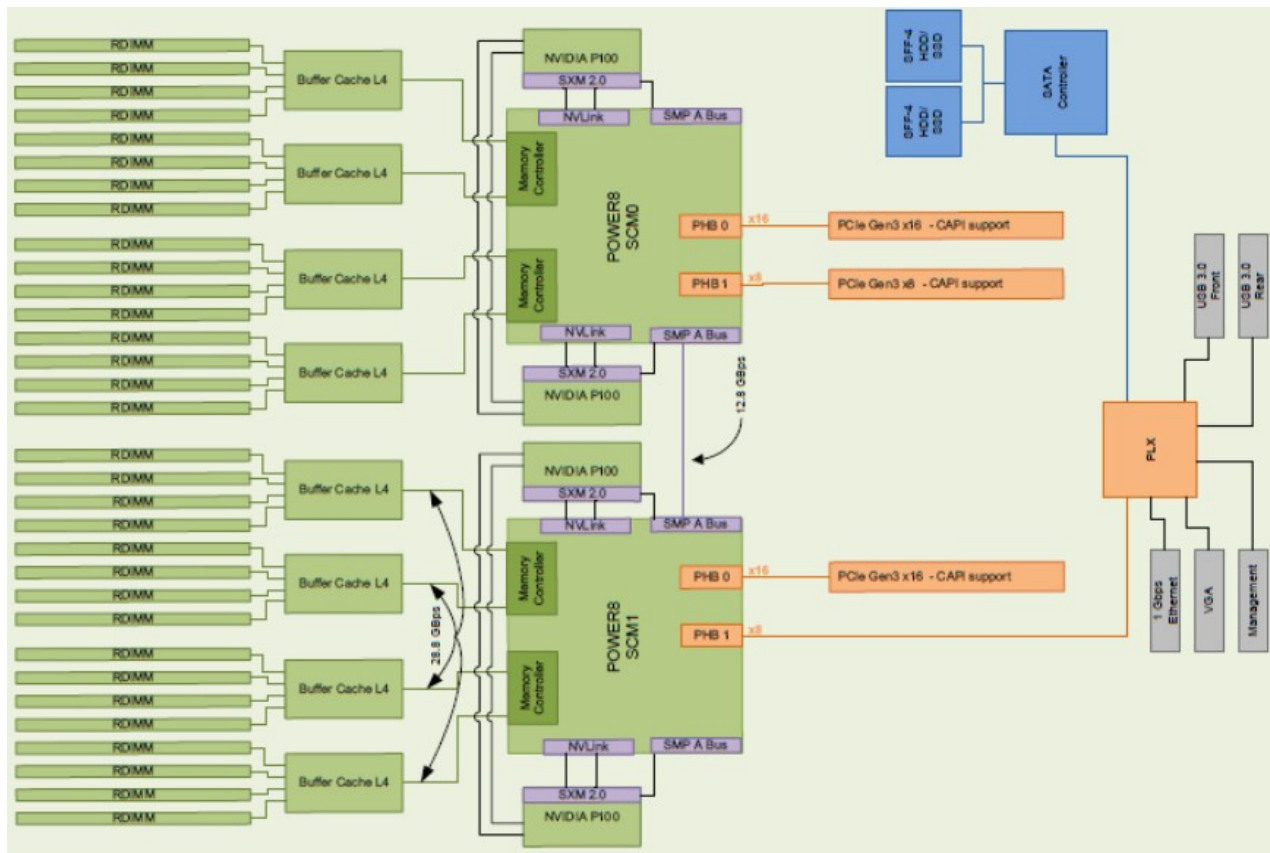
IBM Power System S822LC 'Minsky' (8335-GTB)

Sockets	2 x POWER8
Physical Cores	20
Hardware Threads (Logical Cores)	20 [SMT Off] 40 [SMT 2] 80 [SMT 4] 160 [SMT 8]
CPU Frequency	2.86 GHz [Nominal] 4.03 GHz [Turbo]
Memory Capacity	Up To 1 TB
Memory Bandwidth (Peak)	230 GB/s
DP Performance (Peak)	468 Gflops
GPUs	Up To 4 x NVIDIA Tesla P100
Link	NVLink
Link Bandwidth	20 GB/s
DP Performance (Peak)	4.9 TFlops



Logical System Diagram

- Two Single-Chip Modules (SCM)
- 4 Memory Riser Cards per SCM
 - Buffer Chips for L4 Cache
 - 4 RDIMM Slots
- Max. Capacity: 32 Memory DIMMs (1024 GB)
- 4 GPU Sockets (300 W Max.)
- 3 PCIe Gen3 Slots
- Dedicated PCI Bus:
 - Integrated SATA Controller: Up to 2 SATA Drives
 - Integrated Ethernet
 - Integrated USB Port

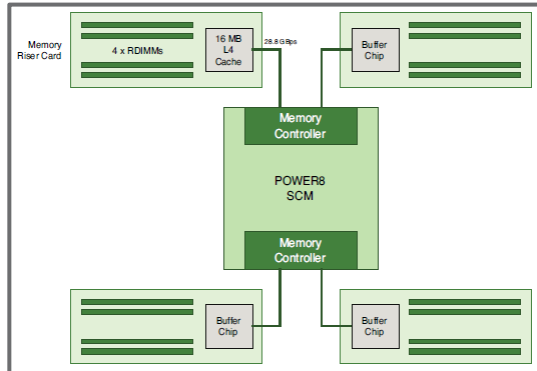
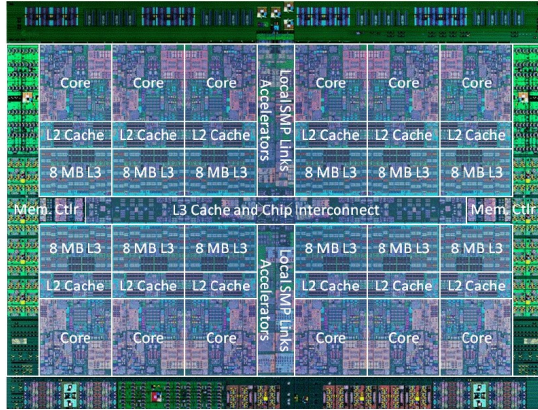


Simultaneous Multi-Threading (SMT)

- Allows a single physical processor core to dispatch simultaneously instructions from more than one hardware thread context
 - With SMT, each POWER8 core can present up to 8 hardware threads
 - Because there are multiple hardware threads per physical processor core, additional instructions can run at the same time
- SMT benefit highly depends on the workload
- Changing SMT mode does not require reboot

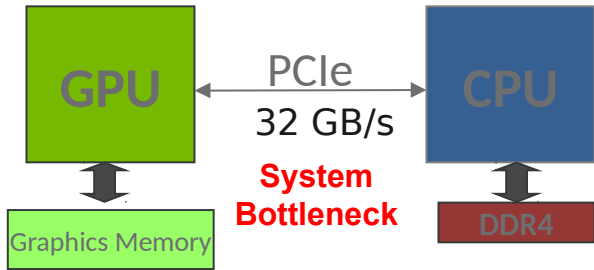
Mode	Logical Cores
Single Thread	20
SMT2	40
SMT4	80
SMT8	160

Memory Sub-System | L3 & L4 Cache

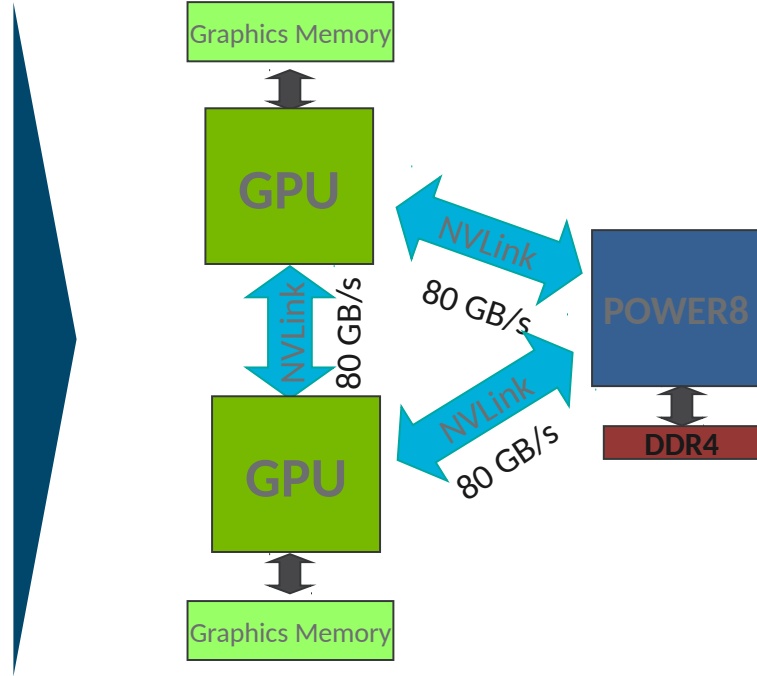


- L3 Cache
 - eDRAM on the processor die
 - Each processor core is associated with a fast 8-MB local region of L3 cache (FLR-L3)
 - But can also access other L3 cache regions as shared L3 cache
- L4 Cache
 - Memory buffer on the memory riser cards
 - Each memory buffer contains 16 MB of L4 cache
 - Up to 128 MB of L4 cache by using all memory riser cards

POWER8 w/NVLink: 2.5x Faster CPU-To-GPU Connection

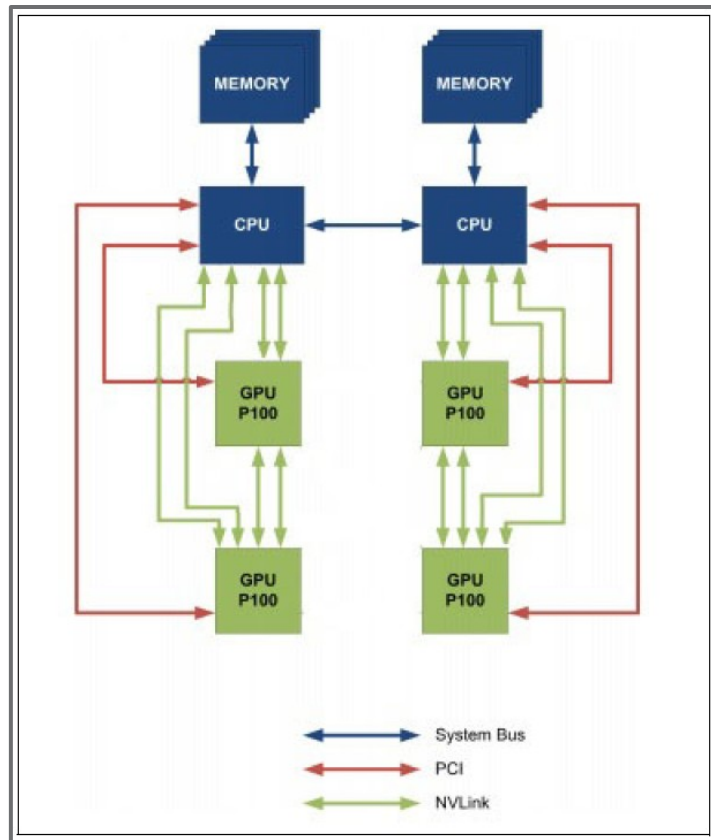
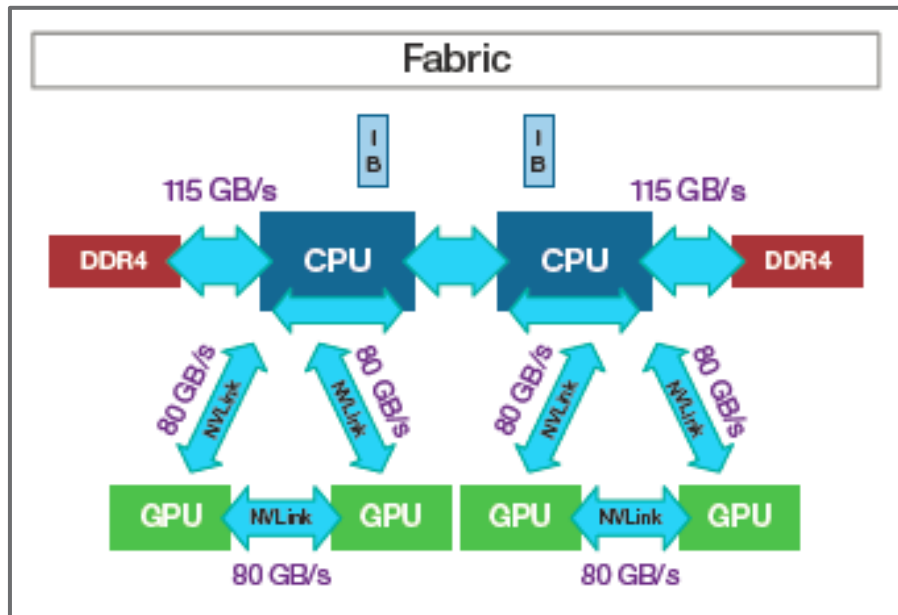


GPUs Limited by PCIe Bandwidth
From CPU-System Memory



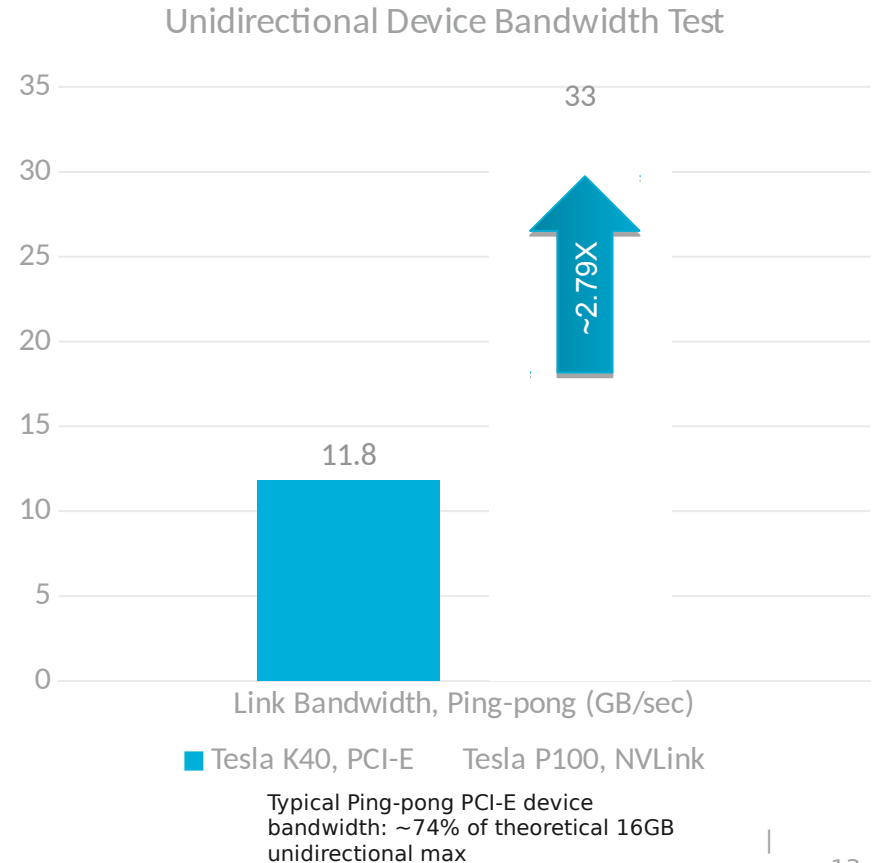
NVLink Enables Fast Unified Memory
Access between CPU & GPU Memories

Comprehensive Fabric Architecture



Unique Bandwidth Dividend: NVLink from CPU:GPU

- **POWER8 with NVLink: the only processor with NVLink from CPU to GPU**
 - Delivering in excess of 2.5x bandwidth in testing *today*
 - NVLink bus delivering higher efficiency than PCI-E links (82.5% vs 74% of peak)
 - No code changes to start leveraging bus (CUDA 8.0 and go)
- **Platform for developers needing Bandwidth for the foreseeable future**
 - POWER8 with NVLink ships in 2016
 - Xeon E5-2600 Series CPUs retain PCI-E x16 3.0 connectivity through 2017



Accelerated Performance: NVIDIA Tesla P100 Pascal GPUs

	Tesla P100	Tesla K80	Tesla K40
DP TFLOPS	5.3 TFLOPS	2.91 TFLOPS (Max Boost)	1.4 TFLOPS
SP TFLOPS	10.6 TFLOPS (21.2 TFLOPS Half Precision)	8.74 TFLOPS (Max Boost)	4.3 TFLOPS
Memory Bandwidth	720 GB/sec	480 GB/sec (2x 240GB)	288 GB/sec
Memory Capacity	16 GB	24 GB (2x 12GB)	12 GB

GIANT LEAPS IN EVERYTHING



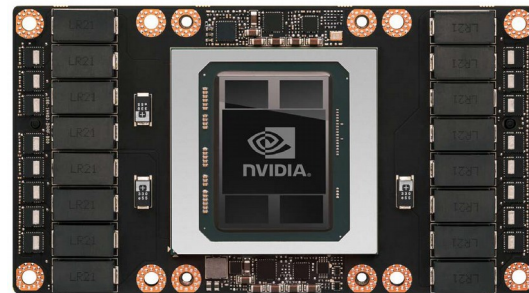
2-3x Compute



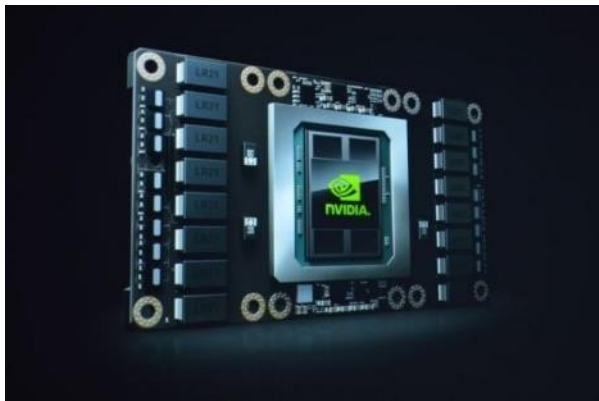
3x GPU Memory BW



5x GPU-GPU BW



NVIDIA Tesla P100 Architecture Details



INTRODUCING TESLA P100

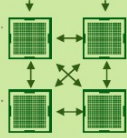
New GPU Architecture to Enable the World's Fastest Compute Node

Pascal Architecture



Highest Compute Performance

NVLink



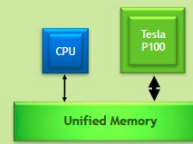
GPU Interconnect for Maximum Scalability

HBM2 Stacked Memory



Unifying Compute & Memory in Single Package

Page Migration Engine



Simple Parallel Programming with 512 TB of Virtual Memory

HPC Software Stack

Software Components

Comprehensive HPC Software Stack

Operating System

- Red Hat Enterprise Linux 7.3 LE

Application Development

- Advance Toolchain 10.0 [Including GCC 6.2]
- CUDA Toolkit 8.0
- GCC 4.8
- IBM XL C/C++ 13.1.5
- IBM XL Fortran 15.1.5
- LLVM Clang & XLFlang [Beta]
- PGI Accelerator 16.10

MPI Libraries

- IBM XL C/C++, IBM XL Fortran:
 - IBM Parallel Environment RTE 2.3 [Phase Out]
 - IBM Spectrum MPI 10.1 [New]
- PGI Accelerator:
 - Open MPI 1.10

Scientific Libraries

- IBM ESSL 5.5
- IBM Parallel ESSL 5.3
- IBM XL MASS

Performance Analysis

- IBM Parallel Performance Toolkit for POWER 2.3

Workload Management

- IBM Spectrum LSF 10.1

Data Management

- IBM Spectrum Scale 4.1

Network Management

- Mellanox Unified Fabric Manager (UFM)

Engineering and Scientific Subroutine Library (ESSL)

ESSL

- Serial & SMP highly-tuned mathematical subroutines
- 30+ SMP CUDA BLAS3 subroutines
 - POWER8 SMP / GPU / Hybrid (CPU+GPU) support
 - Leverage ESSL BLAS + NVIDIA cuBLAS
 - Support multiple GPUs
 - Support problem sizes larger than GPU Memory

Parallel ESSL

- 150+ SPMD highly-tuned mathematical subroutines
 - L2/L3 PBLAS, Linear Algebraic Equations, Fourier Transforms...

Acceleration Enabled Programing Models

**CUDA****Key Features:**

- Gives direct access to the GPU instruction set
- Supports C, C++ and Fortran
- Generally achieves best leverage of GPUs for best application performance
- PGI/NVIDIA Compiler
- CUDA C/C++ for Power via XL NVCC

**Key Features:**

- Designed to simplify Programing of heterogeneous CPU/GPU systems
- Directive based parallelization for accelerator device
- PGI/NVIDIA Compiler
- OpenACC/gcc

**Key Features:**

- OpenMP 4.0 introduces offloading and support for heterogeneous CPU/GPU
- Leverage existing OpenMP high level directives support
- IBM XL Compiler
- Open Source LLVM OpenMP Compiler

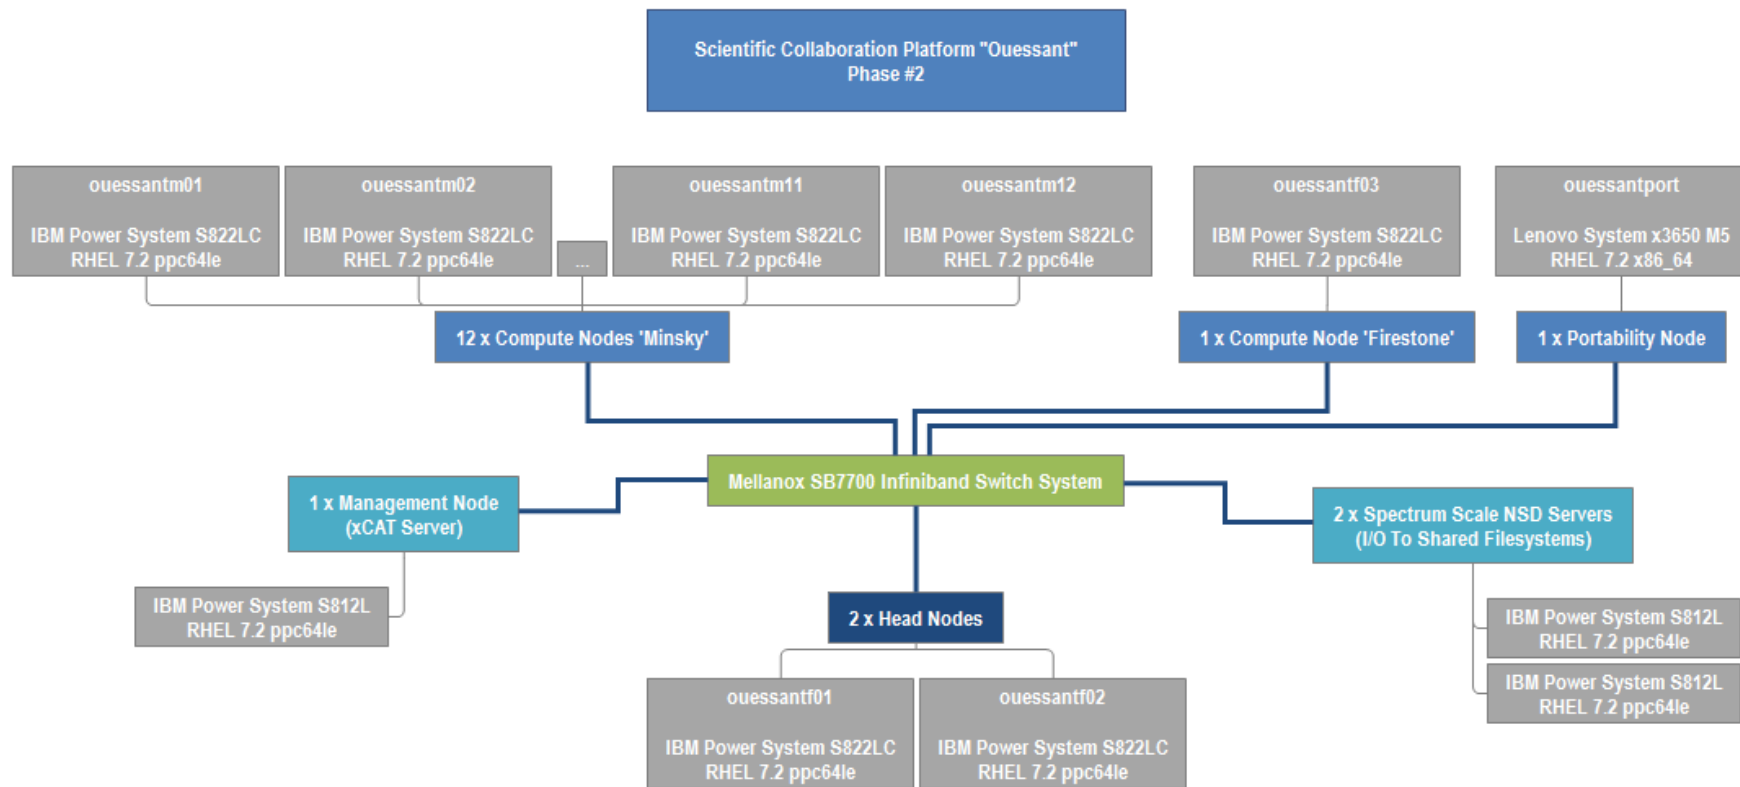
Supported GPU Offloading Features by Compiler Family (1Q 2017)

Compiler	OpenACC	OpenMP
GCC 4	-	4.0 w/o Offload
GCC 5	Experimental 2.0a No Offload on ppc64le	4.0 w/o Offload
GCC 6	Partial 2.0a	Partial 4.0
IBM XL	-	4.0+4.5
LLVM	-	4.0+4.5
PGI	Full 2.0a+2.5	-

Ouessant Platform

Scientific Collaboration Prototype

Platform Architecture



System Topology: CPU

Socket #0												Socket #1											
SMT=0 [Off]																							
0	8	16	24	32	40	48	56	64	72	80	88	96	104	112	120	128	136	144	152				

Socket #0												Socket #1											
SMT=2																							
0	8	16	24	32	40	48	56	64	72	80	88	96	104	112	120	128	136	144	152				
1	9	17	25	33	41	49	57	65	73	81	89	97	105	113	121	129	137	145	153				

Socket #0												Socket #1											
SMT=4																							
0	8	16	24	32	40	48	56	64	72	80	88	96	104	112	120	128	136	144	152				
1	9	17	25	33	41	49	57	65	73	81	89	97	105	113	121	129	137	145	153				
2	10	18	26	34	42	50	58	66	74	82	90	98	106	114	122	130	138	146	154				
3	11	19	27	35	43	51	59	67	75	83	91	99	107	115	123	131	139	147	155				

Socket #0												Socket #1											
SMT=8																							
0	8	16	24	32	40	48	56	64	72	80	88	96	104	112	120	128	136	144	152				
1	9	17	25	33	41	49	57	65	73	81	89	97	105	113	121	129	137	145	153				
2	10	18	26	34	42	50	58	66	74	82	90	98	106	114	122	130	138	146	154				
3	11	19	27	35	43	51	59	67	75	83	91	99	107	115	123	131	139	147	155				
4	12	20	28	36	44	49	60	68	76	84	92	100	108	116	124	132	140	148	156				
5	13	21	29	37	45	50	61	69	77	85	93	101	109	117	125	133	141	149	157				
6	14	22	30	38	46	51	62	70	78	86	94	102	110	118	126	134	142	150	158				
7	15	23	31	39	47	52	63	71	79	87	95	103	111	119	127	135	143	151	159				

System Topology: GPU

- 2 GPU Devices / Socket

- Socket #0

- GPU0

- GPU1

- Socket #1

- GPU2

- GPU3

- “Same Socket” Data Exchanges

- Avoid going through SMP link between POWER8 CPUs

```
[user@host ~]$ nvidia-smi topo --matrix
```

	GPU0	GPU1	GPU2	GPU3	mlx5_0	mlx5_1	CPU Affinity
GPU0	X	NV2	SOC	SOC	SOC	SOC	0-79
GPU1	NV2	X	SOC	SOC	SOC	SOC	0-79
GPU2	SOC	SOC	X	NV2	SOC	SOC	80-159
GPU3	SOC	SOC	NV2	X	SOC	SOC	80-159
mlx5_0	SOC	SOC	SOC	SOC	X	PIX	
mlx5_1	SOC	SOC	SOC	SOC	PIX	X	

Legend:

X = Self
 SOC = Connection traversing PCIe as well as the SMP link between CPU sockets (e.g. QPI)
 PHB = Connection traversing PCIe as well as a PCIe Host Bridge (typically the CPU)
 PXB = Connection traversing multiple PCIe switches (without traversing the PCIe Host Bridge)
 PIX = Connection traversing a single PCIe switch
 NV# = Connection traversing a bonded set of # NVLinks

Job Submission: IBM Spectrum LSF

Typical Submission Script

```
#!/bin/bash
```

```
#BSUB -a p8aff(1,1,1,pack)
```

```
#BSUB -cwd ~/helloworld
```

```
#BSUB -e ~/helloworld/stderr
```

```
#BSUB -J HelloWorld
```

```
#BSUB -n 2
```

```
#BSUB -o ~/helloworld/stdout
```

```
#BSUB -q compute
```

```
#BSUB -R "span[ptile=2]"
```

```
#BSUB -W 00:05
```

```
mpirun ~/helloworld/helloworld.exe
```


Directives

Option	Value	Purpose
-cwd	<Path>	Execution Directory
-e	<File>	stderr File
-J	<Job Name>	Job Name
-n	<# MPI Tasks>	Total Number of MPI Tasks
-o	<File>	stdout File
-q	<Queue>	Target Queue
-R	"span[ptile=<ppn >]"	Resource Specification: Number of Tasks per Node
-W	HH:MM	Run Limit
-x		Travail exclusif

User Commands

Command	Argument	Purpose
bjobs		List active jobs (waiting, in progress)
	-u { <User ID> all }	Restrict to specified user
	-X	List associated resources
bkill	<Job ID>	Kill job
bpeek	<Job ID>	Display job stdout/stderr
	-f	Refresh display in real time
bqueues		List existing queues
bstatus	<Job ID>	Display job status
bsub	< <Submission Script>	Submit job into queue

Execution Environment

Setting	Purpose
Task Placement	Set target compute node for each MPI task (among all allocated hosts)
Processor Affinity	Set list of allowed CPUs for each MPI task
GPU Resource Requirement	Express GPU resources requirement for job submission
GPU Affinity	Set list of allowed GPUs for each MPI task

- **Default Policy: “Group Round Robin”**

- ‘ptile’ MPI tasks per allocated compute node
- One allocated compute node after the other until all MPI tasks have been placed

- **Alternative Policy**

- Specified through environment variable ‘LSB_TASK_GEOMETRY’
- Example:
 - `export LSB_TASK_GEOMETRY="{(0,3)(1,4)(2,5)}«`
 - Tasks #0 & #3 placed on node #1
 - Tasks #1 & #4 placed on node #2
 - Tasks #2 & #5 placed on node #3

- **Purpose**

- Avoid as much as possible resource sharing
- Avoid Linux Scheduler to move processes / threads between CPU cores

- **Management**

- Manual
 - Through LSF Affinity String (Resource Requirement)
- Automated
 - Through LSF Application (esub)
 - New in Spectrum LSF 10.1

- **Tips'n Tricks**

- Checking Required!
 - ALWAYS check applied processor affinity with the help of monitoring tools
 - htop, nmon

- **Application [Automated = Easy Way]**

- Syntax

- #BSUB -a p8aff(num_threads_per_task, SMT, cpus_per_core, distribution_policy)
 - cpus_per_core: # logical CPUs used per physical core
 - distribution_policy = { pack | balance }

- Reference

- <https://goo.gl/VMTtNq>

- **Affinity String [Manual = Hard Way]**

- Syntax

- #BSUB -R affinity[affinity_string]
 - {core|thread}(n):cpubind={core|thread}:distribute={balance|pack}

- Reference

- <https://goo.gl/5v6Qmu>

Processor Affinity: Standard Examples

Execution Configuration	'p8aff' esub Options
2 MPI Tasks x 10 Threads Per Task	#BSUB -a "p8aff(10,1,1,pack)"
4 MPI Tasks x 5 Threads Per Task	#BSUB -a "p8aff(5,1,1,pack)"
10 MPI Tasks x 2 Threads Per Task	#BSUB -a "p8aff(2,1,1,pack)"
20 MPI Tasks x 1 Thread Per Task	#BSUB -a "p8aff(1,1,1,pack)"
20 MPI Tasks x 2 Threads Per Task	#BSUB -a "p8aff(2,2,2,pack)"
20 MPI Tasks x 4 Threads Per Task	#BSUB -a "p8aff(4,4,4,pack)"
20 MPI Tasks x 8 Threads Per Task	#BSUB -a "p8aff(8,8,8,pack)"

GPU Resource Requirement: GPU Compute Mode

- **Available Compute Modes**

- **DEFAULT**

- Equivalent to: Shared

- **EXCLUSIVE_PROCESS**

- Only one process allowed to run

- **EXCLUSIVE_THREAD**

- Only one thread allowed to run

- **PROHIBITED**

- Run not allowed

- **GPU-Specific**

- GPUs can have different Compute Modes

```
[user@host ~]$ nvidia-smi --query --display=COMPUTE
```

```
=====NVSMI LOG=====
```

```
Timestamp                : Tue Nov  1
19:59:57 2016
Driver Version            : 361.93.02

Attached GPUs             : 4
GPU 0002:01:00.0
    Compute Mode          : Default

GPU 0003:01:00.0
    Compute Mode          : Default

GPU 0006:01:00.0
    Compute Mode          : Default

GPU 0007:01:00.0
    Compute Mode          : Default
```

GPU Resource Requirement: Multi-Process Service (MPS)

▪ MPS Benefits

▪ GPU Utilization

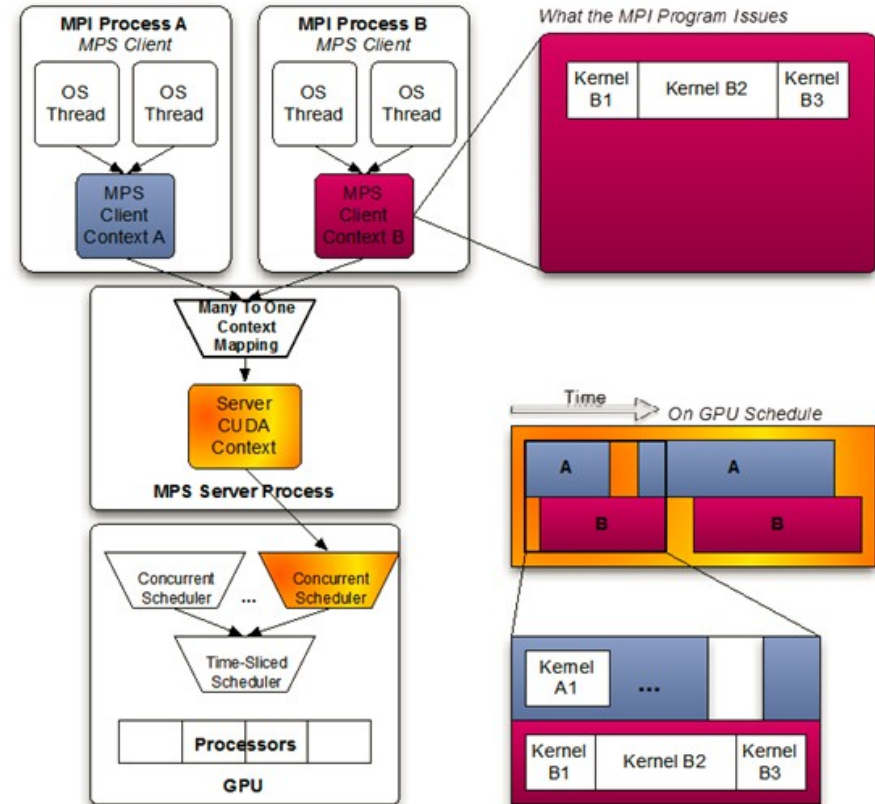
- w/o MPS: Single process may underutilize the compute and memory-bandwidth capacity
- w/MPS: kernel and memcpy operations from different processes can overlap on the GPU, achieving higher utilization and shorter running times

▪ Reduced On-GPU Context Storage

- w/o MPS: each CUDA process using a GPU allocates separate storage and scheduling resources on the GPU
- w/MPS: server allocates one copy of GPU storage and scheduling resources shared by all of its clients

▪ Reduced GPU Context Switching

- w/o MPS: when processes share the GPU their scheduling resources must be swapped on and off the GPU
- w/MPS: server shares one set of scheduling resources between all of its clients, eliminating the overhead of swapping when the GPU is scheduling between those clients



- **Resource Requirement Syntax**

- #BSUB -R "rusage[{ngpus_excl_p|ngpus_shared}=<num_gpus>]"

- **Limitation**

- Spectrum LSF does not manage GPU Compute Mode
 - Statically defined by System Administrator
 - Current state can be examined by the following command
 - lsload -l ngpus:ngpus_excl_p:ngpus_shared <host>

Execution Configuration	Resource Requirement
4 GPUs, Shared Mode	#BSUB -R "rusage[ngpus_shared=2]"
2 GPUs, Exclusive Mode	#BSUB -R "rusage[ngpus_excl_p=2]"
4 GPUs, Exclusive Mode	#BSUB -R "rusage[ngpus_excl_p=4]"

- **Purpose**

- Avoid as much as possible resource sharing
- Prefer GPU direct access (Avoid going through CPU-To-CPU link)

- **Management**

- Manual
 - Through environment variable setting (CUDA_VISIBLE_DEVICES)
- Automated
 - Forthcoming... hopefully!

- **Tips'n Tricks**

- Checking Required!
 - ALWAYS check applied processor affinity with the help of monitoring tools
 - gpustat, nvidia-smi pmon

- **Objective**

- Set CUDA_VISIBLE_DEVICES environment variable with proper value
 - Default value set by Spectrum LSF must be overridden
- Variable should have distinct values for each MPI task
 - At least for exclusive GPU access

- **Suggested Solution**

- Initialize CUDA_VISIBLE_DEVICES value based on MPI rank
- Actual logic might depend on execution configuration
 - # MPI Tasks, Processor Affinity

- **Warning**

- CUDA_VISIBLE_DEVICES="" (null value)
 - Means no GPU assigned to the task
 - Triggers error message at first GPU access

End of Presentation

End Of Presentation