
Jacobi

Project 2016-17


Salvatore Orlando

Regular computation, iterative, stencil-based

Loop until some condition is true

***Perform computation
which involves
communicating with
N,E,W,S neighbors
of a point
(5 point stencil)***

[Convergence test?]



Jacobi

- Jacobi is an iterative method to solve a system of **n linear equations** in **n unknowns** $Ax=b$

- where the i-th equations is the following:

$$a_{i,1}x_1 + a_{i,1}x_1 + \cdots + a_{i,n}x_n = b_i$$

- alternatively, we can say that the i-th equations is:

$$x_i = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j \right]$$

- The various **a**'s and **b**'s are constants, and we have to find the unknowns **x**

Jacobi

- The Jacobi method to solve the linear system is iterative, till we arrive at convergence:

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

where the values of the unknowns at the **(k-1)-th** iteration are used to compute the values at iteration **k**

- The method converges for specific classes of matrixes
 - Usually the solution is approximated till a desired error threshold by computing the difference between the values of the unknowns at the **(k-1)-th** and **k-th** iterations

Jacobi

- The linear system $Ax=b$ can be easily and quickly solved with Jacobi when A is **diagonally dominant**
 - i.e., the magnitude of the diagonal entry in a row is larger than or equal to the sum of the magnitudes of all the other (non-diagonal) entries in that row

$$x_i^k = \frac{1}{a_{i,i}} \left[b_i - \sum_{j \neq i} a_{i,j} x_j^{k-1} \right]$$

Laplace and Jacobi

- Jacobi can be used to solve the differential equation of Laplace in two variables (2D):

$$\frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} = 0$$

- The equation di Laplace models the *steady state* of a function f defined in a physical 2D space, where f is a given physical quantity
- For example, $f(x,y)$ could represent *heat* as measured over a metal plate
 - Given a metal plate, for which we know the temperature at the edges, what is the temperature distribution inside the plate?

Laplace and Jacobi

- For our purposes, it is enough to find the values of $f(x,y)$ for a suitable 2D discretization of the space
- Let Δ be the uniform distance between the discrete grid along the two Cartesian dimensions
 - If Δ is enough small, we can approximate the 2nd order derivatives with the finite difference method. Using the Taylor series, we can thus approximate the 2nd order derivatives as follows:

$$\frac{\partial^2 f}{\partial x^2} = \frac{1}{\Delta^2} [f(x + \Delta, y) - 2f(x, y) + f(x - \Delta, y)]$$

$$\frac{\partial^2 f}{\partial y^2} = \frac{1}{\Delta^2} [f(x, y + \Delta) - 2f(x, y) + f(x, y - \Delta)]$$

- Substituting both in the Laplace equation:


$$f(x, y) = \frac{[f(x - \Delta, y) + f(x, y - \Delta) + f(x + \Delta, y) + f(x, y + \Delta)]}{4}$$

Laplace and Jacobi

- Note the relations between the values of the functions:

$$f(x, y) = \frac{[f(x - \Delta, y) + f(x, y - \Delta) + f(x + \Delta, y) + f(x, y + \Delta)]}{4}$$

- We can note a *cross stencil* of 4 points, which is the same for each (x,y)



STENCIL: a thin sheet of cardboard, plastic, or metal with a pattern or letters cut out of it, used to produce the cut design on the surface below by the application of ink or paint through the holes.

Laplace and Jacobi

- To solve the Laplace equation we have to find all the values of $f(x,y)$ (steady state) wrt a discrete grid with 2 indexes
- To apply Jacobi:
 - we need to solve a linear system, by finding the values of a 1D vector $x(i)$ of unknowns.
 - we thus convert from $f(x,y)$ to $x(i)$
 - We associate the variables $x(i)$ of our linear system $Ax=b$ with the discrete values of $f(x,y)$ with respect to a **discrete 2D grid of n^2 points**
 - Row-major distribution of the 2D discrete matrix derived from $f(x,y)$:

- We thus have n^2 unknowns $x(i)$ with matrix A of size $n^2 \times n^2$

$$\begin{bmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{bmatrix}$$

Laplace and Jacobi

- Since the vector of unknowns $\underline{x(i)}$ of our linear system $Ax=b$ are distributed row-major over the discrete 2D grid of $f(x,y)$

$$\begin{bmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{bmatrix}$$

then:

$$f(x, y) = \frac{f(x - \Delta, y) + f(x, y - \Delta) + f(x + \Delta, y) + f(x, y + \Delta)}{4}$$

becomes:

$$x_i = \frac{x_{i-n} + x_{i-1} + x_{i+1} + x_{i+n}}{4}$$

Laplace and Jacobi

- We have to find matrix **A** and vector **b** of the linear system, on which we have to apply Jacobi
- Our solution to the Laplace equation determined a linear system of equations of this form:

$$x_i = \frac{[x_{i-n} + x_{i-1} + x_{i+1} + x_{i+n}]}{4}$$

that we can re-write as follows:

$$-x_{i-n} - x_{i-1} + 4x_i - x_{i+1} - x_{i+n} = 0$$

- Then, the various $b(i)$ are 0
- and matrix A?

Laplace and Jacobi

- Matrix A of the linear system is **sparse** and is **diagonally dominant**
 - Each row has at most 5 non-zero values
 - All the values on the diagonal are equal to 4, that is $>$ (or \geq) than the sum of all the values on the row

$$\begin{bmatrix} a_{1,1} & \cdots & a_{1,n} \\ \vdots & \ddots & \vdots \\ a_{n,1} & \cdots & a_{n,n} \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} 0 \\ \vdots \\ 0 \end{bmatrix}$$

$$N=9 \Rightarrow n=3$$

$$A = \begin{bmatrix} 4 & -1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ -1 & 4 & -1 & 0 & -1 & 0 & 0 & 0 & 0 \\ 0 & -1 & 4 & -1 & 0 & -1 & 0 & 0 & 0 \\ -1 & 0 & -1 & 4 & -1 & 0 & -1 & 0 & 0 \\ 0 & -1 & 0 & -1 & 4 & -1 & 0 & -1 & 0 \\ 0 & 0 & -1 & 0 & -1 & 4 & -1 & 0 & -1 \\ 0 & 0 & 0 & -1 & 0 & -1 & 4 & -1 & 0 \\ 0 & 0 & 0 & 0 & -1 & 0 & -1 & 4 & -1 \\ 0 & 0 & 0 & 0 & 0 & -1 & 0 & -1 & 4 \end{bmatrix}$$

Jacobi to solve the generated linear system


- Initialize the values of the unknowns $x(i)$, e.g., $x(i) = b(i)$
- The new $x(i)$ are computed iteratively by the Jacobi method as follows:

$$x_i^t = \frac{\left[x_{i-n}^{t-1} + x_{i-1}^{t-1} + x_{i+1}^{t-1} + x_{i+n}^{t-1} \right]}{4}$$

- The new updated values of $x(i)$ at iteration t are used at iteration $t+1$
- Termination condition:

$$\left| x_i^t - x_i^{t-1} \right| < \text{error threshold}(\forall x_i)$$

Example: Heat diffusion with finite differences



Pseudo-codice of Message Passing Jacobi

[Initialize boundary regions, read-only boundary condition]

```
for (i=1; i<=N; i++)  
    x[0][i] = north[i];  
    x[N+1][i] = south[i];  
    x[i][0] = west[i];  
    x[i][N+1] = east[i];
```

[Initialize matrix]

```
for (i=1; i<=N; i++)  
    for (j=1; j<=N; j++)  
        x[i][j] = initvalue;
```

[Iterative refinement of x until values converge]

```
while (maxdiff > CONVERG)
```

[Update x array]


```
for (i=1; i<=N; i++)  
    for (j=1; j<=N; j++)  
        newx[i][j] = 1/4 (x[i-1][j] + x[i][j+1] + x[i+1][j] + x[i][j-1]);
```

[Convergence test]

```
maxdiff = 0;  
for (i=1; i<=N; i++)  
    for (j=1; j<=N; j++)  
        maxdiff = max(maxdiff, |newx[i][j]-x[i][j]|);  
        x[i][j] = newx[i][j]
```



Example

- If we know the boundary conditions of $f(x,y)$, i.e., we know the values of $f(x,y)$ on the edges of the grid and they are constant
 - We can compute the steady state by applying Laplace, and use Jacoby to compute the values of $f(x,y)$ for all the internal points
- Heat diffusion on a metal plate, where the edge of the plate = 1
- Border conditions:
 - $f(x, 0) = 300$
 - $f(x, 1) = f(0, y) = f(1, y) = 0$
- Discretize using a uniform mesh where $\Delta = 0.25$, we obtain the grid on the right-hand




Example

- Once Jacoby solved the Laplace equation, for the discrete points of the 2D mesh, we can plot some iso-temperature curves as follow
 - Remember that the temperature at the bottom edge is kept constant at a value of 300, while is 0 on all the other three edges
- The *isocurves* that connect all the points at the same temperature t are smoother for denser grids, but their computation is more complex in time and space



MPI parallelization: Block data distribution

- Given n^2 points and p^2 processors
 - Every processor is given n^2/p^2 points
 - The four borders, each of size n/p need to be exchanged with the neighbors
 - These four borders are called *ghost area*
 - The communication volume is proportional to $4 * n/p$
 - The computational cost is proportional to n^2/p^2



- What if we increase granularity ?



Heat diffusion: SPMD with block data distribution

- Matrix of n^2 elements distributed (block,block) on a grid of p^2 processes
- Each process $P_{i,j}$ is thus assigned a block of n^2/p^2 elements
- Code executed by a generic SPMD process $P_{i,j}$

```
<declaration of array h(0:n/p+1, 0:n/p+1) with ghost area>
<declaration of array temp. g(1:n/p, 1:n/p) >
for (iteration = 0; iteration < limit; iteration++) {
    < border exchange of h[][] with the 4 neighbors  $P_{i,j-1}$   $P_{i,j+1}$   $P_{i-1,j}$   $P_{i+1,j}$  > | Local
                                                                    | Barrier
    for (i = 1; i <= n/p; i++)
        for (j = 1; j <= n/p; j++)
            g[i][j] = 0.25*(h[i-1][j]+h[i+1][j]+h[i][j-1]+h[i][j+1]);
    for (i = 1; i <= n/p; i++) /* Update h[][] */
        for (j = 1; j <= n/p; j++)
            h[i][j] = g[i][j];
}
```



- We have to differentiate the code for blocks/processes on the grid borders
 - $P_{i,j}$ where $i, j=0, p-1$
 - less partners with which to exchange communications
- With asynchronous communications, we can postpone the receives
 - the receive must however be completed before accessing the ghost area

Increasing granularity



- 64 points
- 64 computations
- 4 points sent per processor
- $4 * n_{\text{procs}} = 4 * 64 =$ communication volume 256

- 64 points
- 64 computations
- 16 points sent per processor
- $16 * n_{\text{procs}} = 16 * 4 =$ communication volume 64




Surface / Volume

- The **communication cost** is proportional to the surface (perimeter) of the data partition
- The **computational cost** is proportional to the volume (area) of the data partition
- In two dimensions:
 - The surface increases as n
 - The volume increases as n^2
- In three dimensions:
 - The surface increase as n^2
 - The volume increases as n^3
- The ratio **communication/computation** decreases when increasing the partition size
 - Good !
 - This effect is less visible for high dimensional problems

MPI parallelization: Stripe data distribution

- Given n^2 points and p^2 processors
 - Every processor is given n^2/p^2 points
 - Those points form a $n * n/p^2$ rectangle
 - Two borders, each of size n need to be exchanged with the neighbors
 - The data exchanged are proportional to n (2 sends and 2 receives)
- In the (block,block) distribution
 - The data exchanged are proportional to n/p (4 sends and 4 receives)
- To minimize the communication volume:
 - (block,block) is better




Project


- Design Jacobi in parallel and evaluate performance plots
 - Speedup, the same problem size, ...
 - Scalability, how the execution time gets larger as we increases the size of the problem (e.g., by reducing Δ)
- MPI parallelization
- Choose some of these
 - OpenMP parallelization
 - OpenMP & MPI parallelization
 - SIMD parallelization
 - GPU parallelization

Some issues in implementing Jacobi


- Synchronization
 - Can we avoid to synchronize at each iteration?
 - Measure the effect on convergence, and if this strategy save time
- Message passing (MPI) implementation
 - Can we avoid to update the borders owned by other tasks at each iteration?
Effect on convergence and execution time.
- Data partitioning
 - Block or strip?



Block



Uniform Strip



Non-uniform Strip